

# CNN-based Disease Prediction for Crops

---

Shreyas Anil, Ryan Bernstein, Aryan Mistry, Justin Yang


# The Problem

- Globally, more than 20% of crop yields are lost to plant disease every year.
- Plant disease contributes to food insecurity, poverty, and famine, especially in lesser-developed countries.
- There is a pressing need for a reliable early-detection system, which would allow farmers to save blighted crops.

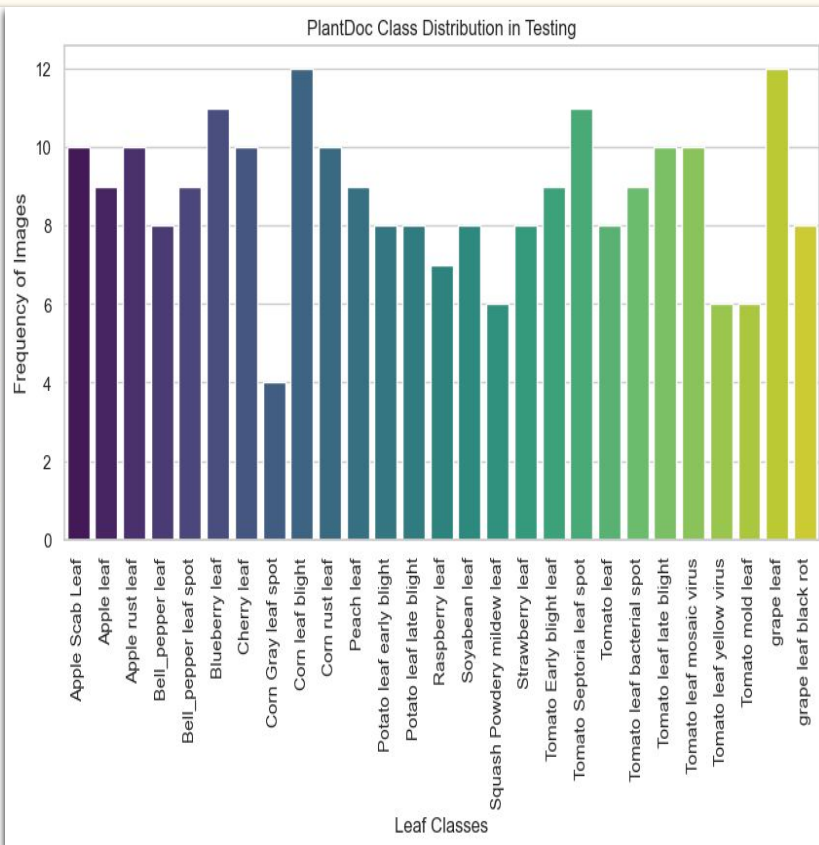



# Objectives

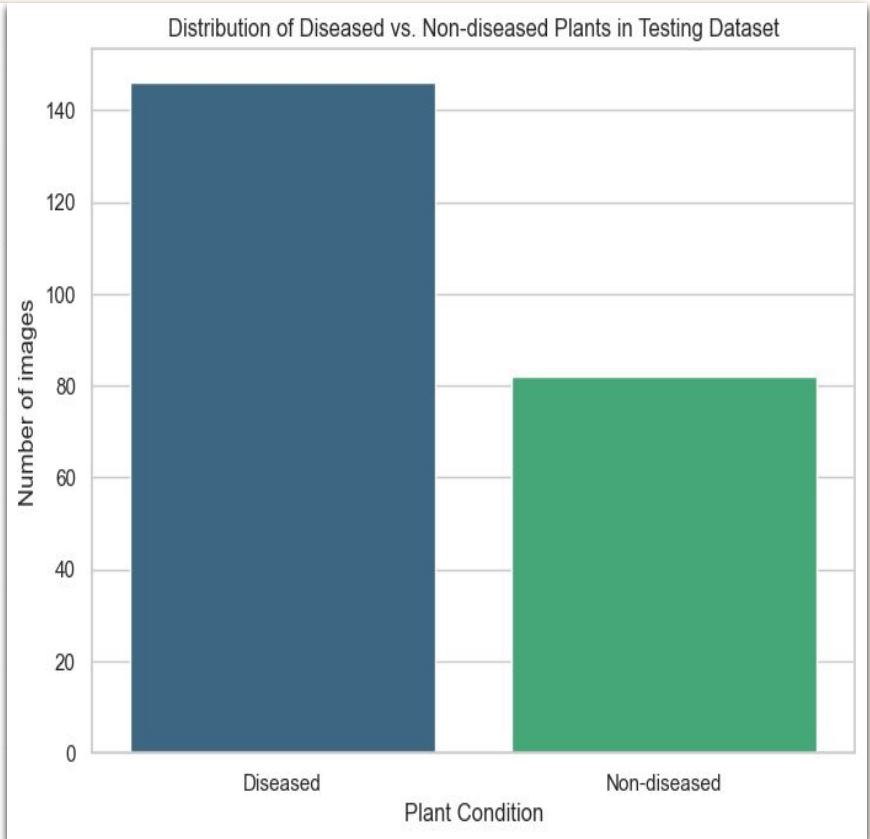
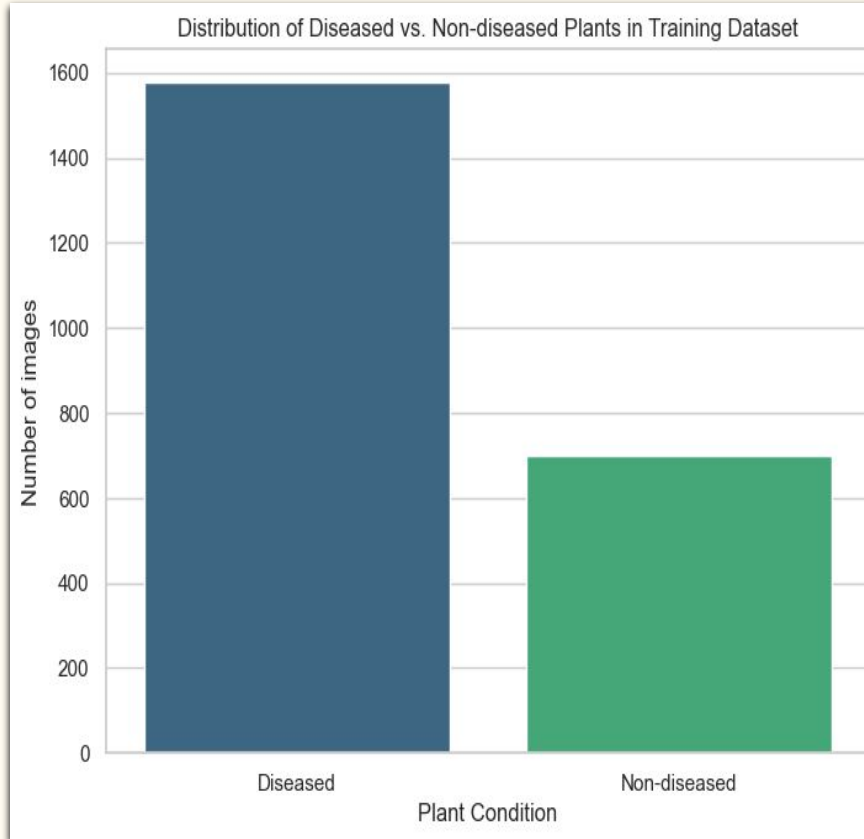
- We aim to implement a computer-vision deep learning model that classifies a plant image by species, and predicts whether it is diseased or not.
- We leverage the PlantDoc dataset (~2600 images from 13 different species and 17 different disease types) for model training.



# Data Structure - Species



# Data Structure - Diseases



# Methodology

## **Data Preprocessing/Augmentation:**


1. Resized and cropped input images
2. Randomly introduced horizontally flipped images into training data
3. Converted images to tensors and normalized color via ImageNet
4. Used a DataLoader with reshuffling for developing batches

## **Modeling**

- A Simple CNN with 9 layers
  - Vanilla
  - w/Regularization
  - w/Stochastic Gradient Descent, Regularization and Scheduler
- ResNet 18 and 50
- EfficientNet 0 and B4
- All trained for only 10 epochs

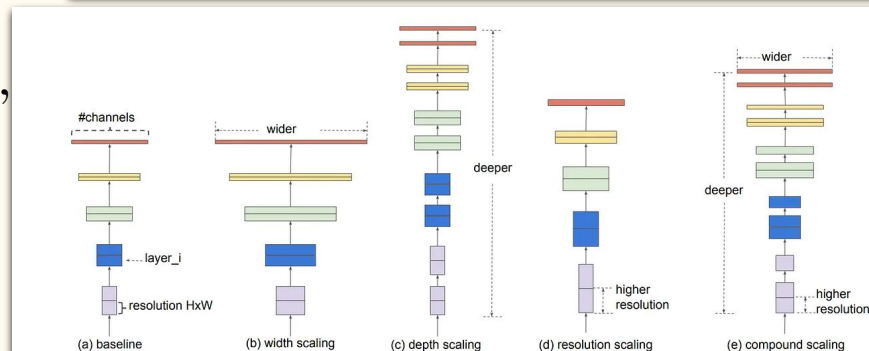

# ResNet

- Developed by Microsoft Research
- Builds on existing CNNs with a New Layer: Residual Blocks
  - Contains “Skip Connection”
- We utilized ResNets already pretrained on ImageNet with 18 and 50 layers
- Helps solve the Vanishing/Exploding Gradient Problem, while also easily adaptable



# EfficientNet

- Developed by Google AI
- Employs **compound scaling**, which uniformly scales all dimensions of depth, width, and resolution
- We utilized EfficientNet B0 and B4, as depicted in the graph above
  - Pretrained on ImageNet
- Helps scale parameters simultaneously, allowing for better utilization of computational resources





# Code Overview

```
[ ] res_50_model = create_pretrained_model(num_classes)

# Set up loss and optimizer
criterion = nn.CrossEntropyLoss()
optimizer = torch.optim.Adam(res_50_model.parameters(), lr=0.001)

# Training and testing
device = torch.device("cuda" if torch.cuda.is_available() else "cpu")
```

➔ /usr/local/lib/python3.10/dist-packages/torchvision/models/\_utils.py:208: UserWarning: The parameter 'pretrained' is deprecated since 0.13 and may be removed in warnings.warn(  
/usr/local/lib/python3.10/dist-packages/torchvision/models/\_utils.py:223: UserWarning: Arguments other than a weight enum or `None` for 'weights' are deprecated warnings.warn(msg)  
Downloading: "<https://download.pytorch.org/models/resnet50-0676ba61.pth>" to /root/.cache/torch/hub/checkpoints/resnet50-0676ba61.pth  
100%|██████████| 97.8M/97.8M [00:00<00:00, 113MB/s]

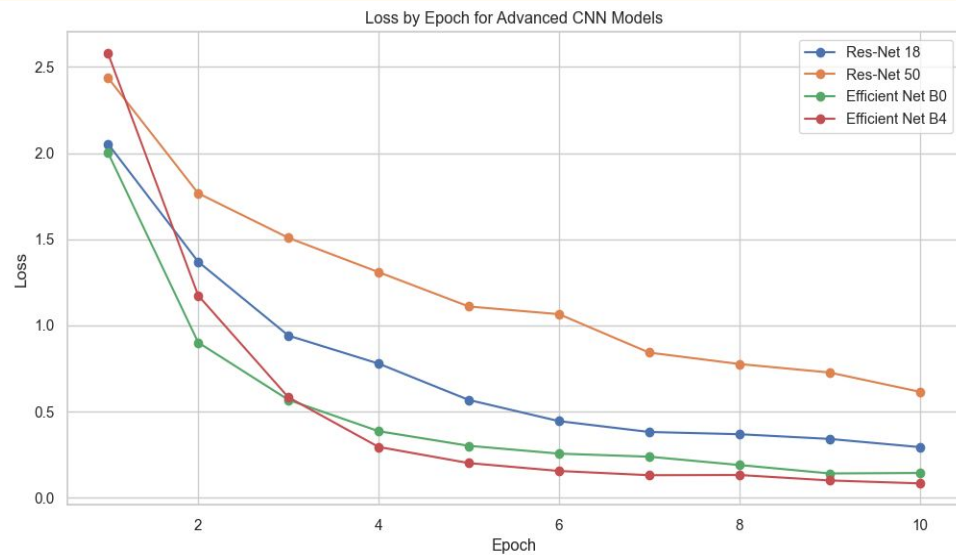
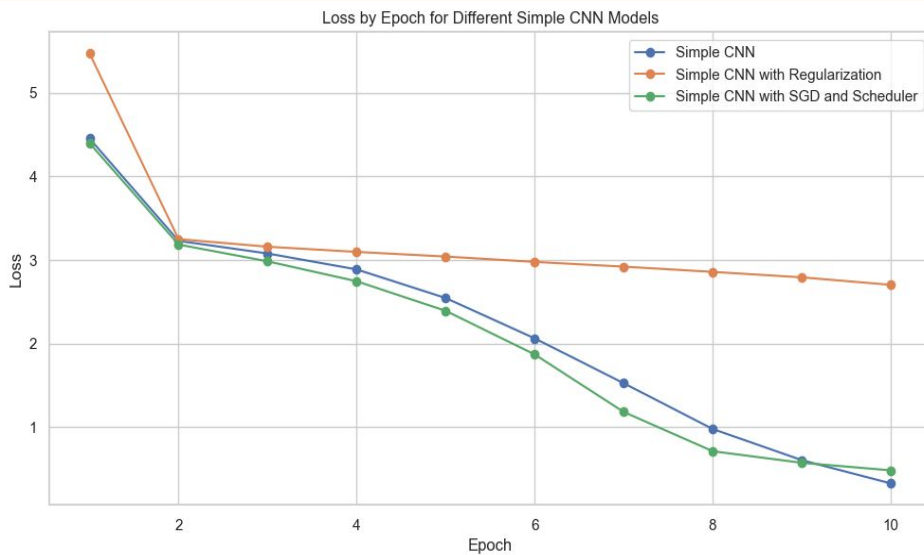
```
[ ] train_model(res_50_model, train_loader, criterion, optimizer, num_epochs=10, device=device)
```

➔ Epoch 1, Loss: 2.43699514543688  
Epoch 2, Loss: 1.7665844930184853  
Epoch 3, Loss: 1.5082206436105676  
Epoch 4, Loss: 1.3089014195107125  
Epoch 5, Loss: 1.1097415427903872  
Epoch 6, Loss: 1.0651415248174925  
Epoch 7, Loss: 0.8421571931323489  
Epoch 8, Loss: 0.7758747323139293  
Epoch 9, Loss: 0.7265614617515255  
Epoch 10, Loss: 0.6148334806029861

```
[ ] test_model(res_50_model, test_loader, device=device)
```

➔ Accuracy of the network on the test images: 39.40677966101695%  
F1 Score (Weighted): 0.37661299605204734

# Results - Model Losses over Epochs



# Results-Accuracy Table

Accuracy and F1 Score for Simple CNNs

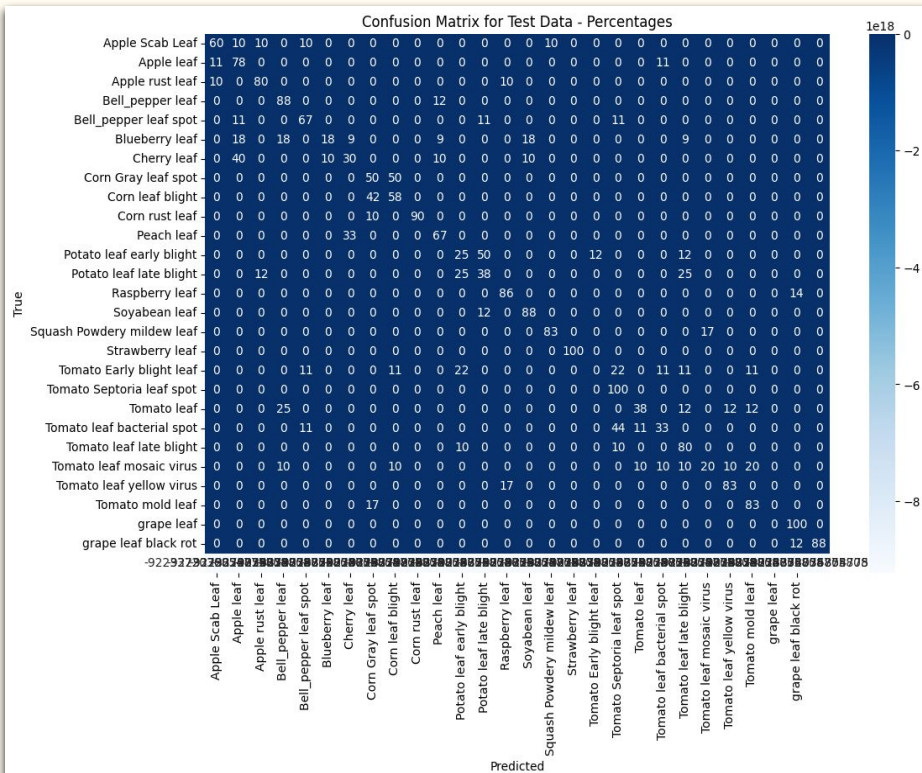

Model	Accuracy (%)	F1 Score (Weighted)
Simple CNN	18.220338983050848	0.15146084728608208
Simple CNN with Regularization	11.016949152542374	0.08775200407716652
Simple CNN with SGD and Scheduler	16.10169491525424	0.1480259592552006

Accuracy and F1 Score for Advanced CNNs

Model	Accuracy (%)	F1 Score (Weighted)
Res-Net 18	40.25423728813559	0.3666380717501854
Res-Net 50	39.40677966101695	0.37661299605204734
Efficient Net B0	47.45762711864407	0.47404671235547036
Efficient Net B4	52.54237288135593	0.5177276451181709



## Results - Confusion Matrix for EfficientNet B4



# Conclusion and Next Steps

- EfficientNet B4 was our best model; it achieved impressive accuracy despite only being trained for 10 epochs.
- Complex CNNs with pre-trained weights perform significantly better than simple CNNs without pre-trained weights
- In the future we may want to look into doing bounding box implementation
- We could also do a binary classification task to see how well a model can distinguish between diseased plants and non-diseased plants

# Sources and Extra Links

[1] Savary, S., Willocquet, L., Pethybridge, S.J. *et al.* The global burden of pathogens and pests on major food crops. *Nat Ecol Evol* 3, 430–439 (2019).

<https://doi.org/10.1038/s41559-018-0793-y>

[2] Singh, Davinder, et al. "PlantDoc: A dataset for visual plant disease detection." *Proceedings of the 7th ACM IKDD CoDS and 25th COMAD*. 2020. 249-253.

[3] He, Kaiming, et al. "Deep residual learning for image recognition." *Proceedings of the IEEE conference on computer vision and pattern recognition*. 2016.

[4] Tan, Mingxing, and Quoc Le. "Efficientnet: Rethinking model scaling for convolutional neural networks." *International conference on machine learning*. PMLR, 2019.

Link to GitHub: [https://github.com/jyang0620/Plant\\_Disease\\_Classifier](https://github.com/jyang0620/Plant_Disease_Classifier)

- Medium article coming soon, will be linked in the GitHub